

# Analysis of 77072 zip code

## Imports and loading csv

```
In [141]: #Imports
import pandas as pd
import numpy as np
from pandas.plotting import register_matplotlib_converters
import matplotlib.pyplot as plt
from matplotlib.pylab import rcParams
register_matplotlib_converters()

from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

from scipy import stats
from random import gauss as gs
import datetime

from statsmodels.tsa.arima_model import ARMA
from statsmodels.tsa.stattools import adfuller, acf, pacf
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import statsmodels.api as sm
from statsmodels.tsa.seasonal import seasonal_decompose

#Supress default INFO logging
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
import logging
logger = logging.getLogger()
logger.setLevel(logging.CRITICAL)
import logging, sys
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
In [142]: df = pd.read_csv('Data Files/df_zillow_77072_prepped_fbprophet.csv')
```

```
In [143]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 265 entries, 0 to 264
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    ds      265 non-null    object
1    y        265 non-null    float64
dtypes: float64(1), object(1)
memory usage: 4.3+ KB
```

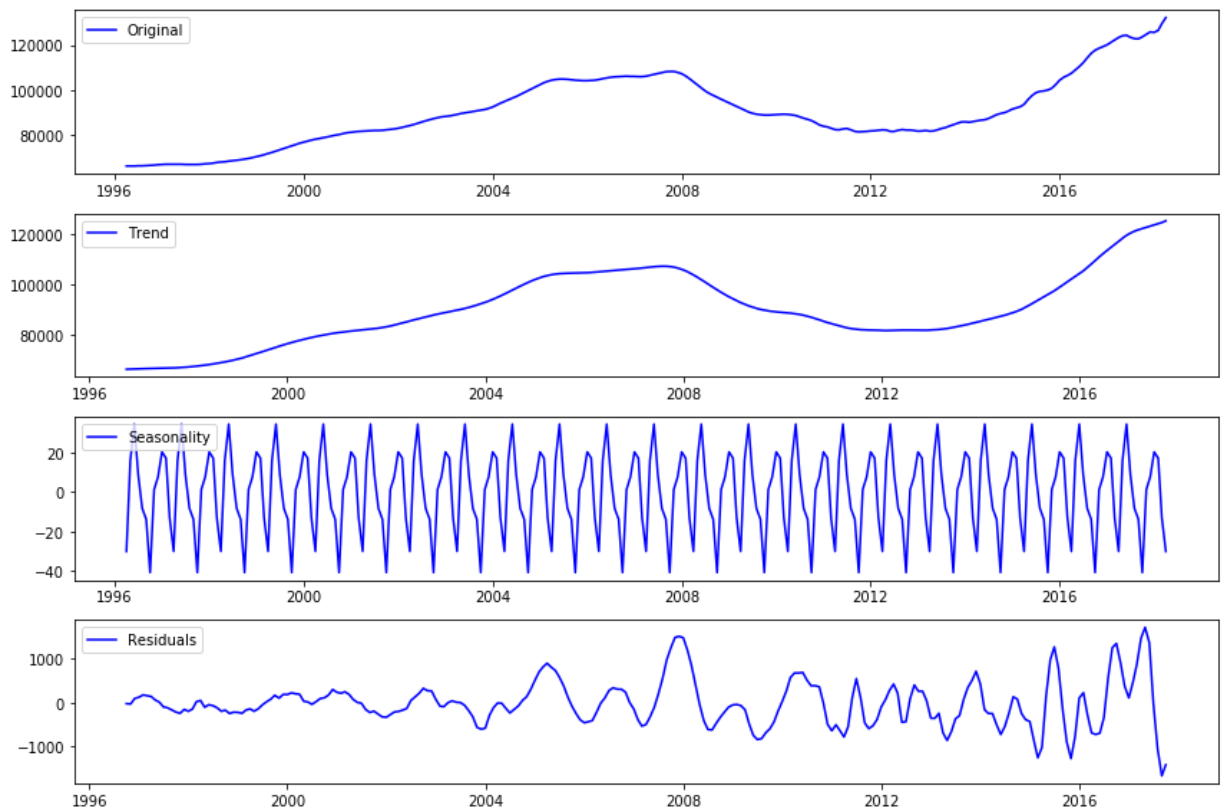
## Decomposition and plots

```
In [144]: df.index = pd.to_datetime(df['ds'])  
df= df.drop(columns='ds')
```

```
In [145]: decomposition = seasonal_decompose(df.y)  
observed = decomposition.observed  
trend = decomposition.trend  
seasonal = decomposition.seasonal  
residual = decomposition.resid
```

```
In [146]: register_matplotlib_converters()
```

```
In [147]: plt.figure(figsize=(12,8))  
plt.subplot(411)  
plt.plot(observed, label='Original', color="blue")  
plt.legend(loc='upper left')  
plt.subplot(412)  
plt.plot(trend, label='Trend', color="blue")  
plt.legend(loc='upper left')  
plt.subplot(413)  
plt.plot(seasonal, label='Seasonality', color="blue")  
plt.legend(loc='upper left')  
plt.subplot(414)  
plt.plot(residual, label='Residuals', color="blue")  
plt.legend(loc='upper left')  
plt.tight_layout()
```



**I want to see if the data correlates with earlier data of**

## itself

1) Get rolling average with window of 4

- Couldn't see much with window of 1-3


2) Plot data against itself with rolling avg to see visual of the graph.

```
In [148]: df['roll_avg'] = df.rolling(window=4).mean()  
df.corr()
```

Out[148]:

	y	roll_avg
y	1.000000	0.998385
roll_avg	0.998385	1.000000


```
In [149]: plt.figure(figsize=(20, 10))  
plt.scatter(df.index[:265], df['y'][:265], s=20)  
plt.scatter(df.index[1:265], df['roll_avg'][1:265], s=20);
```



```
In [150]: lr = LinearRegression()  
lr.fit(df[['roll_avg']][4:], df['y'][4:])
```

Out[150]: LinearRegression(copy\_X=True, fit\_intercept=True, n\_jobs=None, normalize=False)

```
In [151]: plt.figure(figsize=(20, 10))
plt.plot(df.index[:265], df['y'][:265], label='Data')
plt.plot(df.index[4:265], lr.predict(df[['roll_avg']][4:265]),
         label='Predicted')
plt.legend();
```



**Upon brief visual look, there might be some correlation. We will set up for our model by using the Dickey-Fuller test and ACF (Auto-correlation) and PACF (Partial-autocorrelation)**

## Checking for Stationarity

```
In [152]: dfctest = adfuller(df.y)
dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value', '#Lags Used']
for key,value in dfctest[4].items():
    dfoutput['Critical Value (%)'%key] = value
print(dfctest)
print()
print(dfoutput)
```

```
(-0.7060649545040951, 0.8451723171119757, 12, 252, {'1%': -3.4565688966099373,
'5%': -2.8730786194395455, '10%': -2.5729189953388762}, 3452.629967349372)
```

```
Test Statistic          -0.706065
p-value                 0.845172
#Lags Used              12.000000
Number of Observations Used  252.000000
Critical Value (1%)      -3.456569
Critical Value (5%)      -2.873079
Critical Value (10%)     -2.572919
dtype: float64
```

### Dickey Fuller Test

- We see that test statistic value is -0.706065
- We see that the critical values are LESS than the test statistic. (-3.45, -2.87, -2.57)
- From just the baseline data, the test statistic I have is MORE than the critical value.
- We accept the null that the time series is not stationary!



### P-Value analysis

1. If p-value > 0.05: Fail to reject the null hypothesis (H0), the data has a unit root and is non-stationary.
  - Our current p-value is 0.845172
    - This means: p-value > 0.05: Fail to reject the null hypothesis (H0), the data has a unit root and is non-stationary.
2. If p-value <= 0.05: Reject the null hypothesis (H0), the data does not have a unit root and is stationary.
  - Our goal is to make the data stationary

## Auto-Correlation and Partial Auto-Correlation Check


```
In [153]: #ACF using plotting
plt.figure(figsize=(20, 5))
pd.plotting.autocorrelation_plot(df['y']);

#Statsmodels ACF
rcParams['figure.figsize'] = 20, 5
plot_acf(df['y'], lags=125, alpha=0.05);
```



**PACF**

```
In [154]: pacf(df['y'], nlags=20)
rcParams['figure.figsize'] = 20, 5
plot_pacf(df['y'], lags=50, alpha=0.05);
```



## Observations of ACF and PACF

### We see the following:

- We know that the ACF describes the autocorrelation between an observation and another observation at a prior time step that includes direct and indirect dependence information.
  - After about 17-18 lags, the line goes into our confidence interval (light blue area).
  - This can be due to seasonality of every 18 months in our data.
- We know that the PACF only describes the direct relationship between an observation and its lag.
  - PACF cuts off after lags = 2
  - This means there are no correlations for lags beyond 2

**\*\* Granted the data is not stationary, we will have to transform the data to make it stationary and satisfy the Dicky-Fuller test\*\***

## De-trending and transforming the data


1. I will try the following
  - Log transform
  - Subtract rolling mean

- Run Dickey-Fuller test with each transform to see if I can reject/accept the null hypothesis
- Null-Hypothesis for Dickey-Fuller test is: The null-hypothesis for the test is that the time series is not stationary. So if the test statistic is less than the critical value, we reject the null hypothesis and say that the series is stationary.

## Log-transform on data and testing for stationarity

```
In [155]: logged_df = df['y'].apply(lambda x : np.log(x))
```

```
In [156]: ax1 = plt.subplot(121)
logged_df.plot(figsize=(12,4), color="tab:green", title="Log Transformed Values")
ax2 = plt.subplot(122)
df.plot(color="tab:blue", title="Original Values", ax=ax2);
```



```
In [157]: dfctest = adfuller(logged_df)
dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value', '#Lags Used']
for key,value in dfctest[4].items():
    dfoutput['Critical Value (%)'%key] = value
print(dfctest)
print()
print(dfoutput)
```

```
(-1.1458067631277982, 0.6964635146228962, 12, 252, {'1%': -3.4565688966099373,
'5%': -2.8730786194395455, '10%': -2.5729189953388762}, -2273.0439965308883)
```

Test Statistic	-1.145807
p-value	0.696464
#Lags Used	12.000000
Number of Observations Used	252.000000
Critical Value (1%)	-3.456569
Critical Value (5%)	-2.873079
Critical Value (10%)	-2.572919
dtype:	float64

## Observations after log-transform



1. Test Statistic is still larger than Critical Values. We accept the null-hypothesis that the time series is not stationary!
  - Test Statistic -1.145807
  - Critical Value (1%) -3.456360
  - Critical Value (5%) -2.872987
  - Critical Value (10%) -2.572870
2. P value is 0.696464
  - This means: p-value > 0.05: Fail to reject the null hypothesis (H0), the data has a unit root and is non-stationary.

## Subtracting Rolling Mean from logged data and a better window size

```
In [158]: #Try breakdown with data minus rollmean. It looks like there is seasonality but  
# Window of 11
```

```
logged_df_roll_mean = logged_df.rolling(window=11).mean()  
logged_df_minus_roll_mean1 = logged_df - logged_df_roll_mean  
logged_df_minus_roll_mean1.dropna(inplace=True)
```

```
In [159]: logged_df_minus_roll_mean1.head()
```

```
Out[159]: ds  
1997-02-01    0.007671  
1997-03-01    0.006574  
1997-04-01    0.005477  
1997-05-01    0.004380  
1997-06-01    0.003421  
Name: y, dtype: float64
```

```
In [160]: dfctest = adfuller(logged_df_minus_roll_mean1)
# Extract and display test results in a user friendly manner
dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value', '#Lags Used']
for key,value in dfctest[4].items():
    dfoutput['Critical Value (%s)'%key] = value
print(dfctest)
print()
print(dfoutput)
```

```
(-1.596689364254372, 0.48525631721789897, 12, 242, {'1%': -3.457664132155201,
'5%': -2.8735585105960224, '10%': -2.5731749894132916}, -2216.5553842190684)
```

```
Test Statistic          -1.596689
p-value                  0.485256
#Lags Used               12.000000
Number of Observations Used 242.000000
Critical Value (1%)      -3.457664
Critical Value (5%)      -2.873559
Critical Value (10%)     -2.573175
dtype: float64
```

## --- Observations from Dickey Fuller Test ---

**We are getting close.**

- Test statistic is -1.596689 which is higher than crit values
- p value is 0.485256 , I cannot reject null

## Differencing the data and re-running Dickey Fuller

```
In [161]: logged_df_diff = logged_df.diff(periods=1)
```

```
In [162]: logged_df_diff_roll_mean = logged_df_diff.rolling(window=11).mean()
logged_df_diff_roll_mean1 = logged_df_diff - logged_df_diff_roll_mean
logged_df_diff_roll_mean1.dropna(inplace=True)
```

```
In [163]: logged_df_diff_roll_mean1.head()
```

```
Out[163]: ds
1997-03-01    -0.001097
1997-04-01    -0.001097
1997-05-01    -0.001097
1997-06-01    -0.000959
1997-07-01    -0.002323
Name: y, dtype: float64
```

```
In [164]: dfctest = adfuller(logged_df_diff_roll_mean1)
# Extract and display test results in a user friendly manner
dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value', '#Lags Used']
for key,value in dfctest[4].items():
    dfoutput['Critical Value (%)'%key] = value
print(dfctest)
print()
print(dfoutput)
```

```
(-4.556078938977154, 0.0001555133374368479, 11, 242, {'1%': -3.457664132155201,
'5%': -2.8735585105960224, '10%': -2.5731749894132916}, -2206.246504084256)
```

Test Statistic	-4.556079
p-value	0.000156
#Lags Used	11.000000
Number of Observations Used	242.000000
Critical Value (1%)	-3.457664
Critical Value (5%)	-2.873559
Critical Value (10%)	-2.573175
dtype:	float64

### --- Observations of Dickey-Fuller Test ---


- We see Test Statistic is less than the Critical values, this satisfies the stationarity assumption. We can reject the null and say series is stationary.

- Test Statistic	-4.556079
- Critical Value (1%)	-3.458247
- Critical Value (5%)	-2.873814
- Critical Value (10%)	-2.573311


- We see that p-value = 0.000156. Since  $p \leq 0.05$ , I can reject the null hypothesis ( $H_0$  = series is non-stationary). The data does not have a unit root and is stationary.

## ACF and PACF

```
In [165]: #Statsmodels ACF
rcParams['figure.figsize'] = 20, 5
plot_acf(logged_df_diff_roll_mean1, lags=125, alpha=0.05);
```



```
In [166]: #PACF plot
rcParams['figure.figsize'] = 20, 4
plot_pacf(logged_df_diff_roll_mean1, lags=50, alpha=0.05);
```



### --- Observations of ACF and PACF ---

1. After about 4 lags, the line goes into our confidence interval (light blue area).
  - This can be due to seasonality of every 4 months in our data.
2. PACF trails off after 3-4 lags.
  - Also slight slight sinusoidal behavior but nothing crazy
  - This means there are no high correlations for lags beyond 2-3
3. Based on above information and that the data is stationary, we can use the p and q values for the ARMA model
  - $p = 4$  (per ACF)
  - $q = 2, 3, 4$  (per PACF)

## ARMA Modeling

```
In [167]: # Instantiate & fit model with statsmodels
#p = num lags - ACF
p = 4

# q = Lagged forecast errors - PACF
q = 4

# Fitting ARMA model and summary
ar = ARMA(logged_df_minus_roll_mean1,(p, q)).fit()
ar.summary()
```

Out[167]: ARMA Model Results

<b>Dep. Variable:</b>	y	<b>No. Observations:</b>	255
<b>Model:</b>	ARMA(4, 4)	<b>Log Likelihood</b>	1196.797
<b>Method:</b>	csm-mle	<b>S.D. of innovations</b>	0.002
<b>Date:</b>	Thu, 29 Apr 2021	<b>AIC</b>	-2373.594
<b>Time:</b>	14:17:07	<b>BIC</b>	-2338.181
<b>Sample:</b>	02-01-1997	<b>HQIC</b>	-2359.349
	- 04-01-2018		

	coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	0.0159	0.011	1.464	0.145	-0.005	0.037
<b>ar.L1.y</b>	2.6469	0.101	26.121	0.000	2.448	2.845
<b>ar.L2.y</b>	-2.4693	0.281	-8.798	0.000	-3.019	-1.919
<b>ar.L3.y</b>	0.7846	0.281	2.797	0.006	0.235	1.334
<b>ar.L4.y</b>	0.0297	0.101	0.294	0.769	-0.168	0.228
<b>ma.L1.y</b>	-0.4997	0.081	-6.133	0.000	-0.659	-0.340
<b>ma.L2.y</b>	-0.5621	0.092	-6.123	0.000	-0.742	-0.382
<b>ma.L3.y</b>	0.1243	0.063	1.968	0.050	0.001	0.248
<b>ma.L4.y</b>	0.6447	0.067	9.687	0.000	0.514	0.775

Roots

	Real	Imaginary	Modulus	Frequency
<b>AR.1</b>	1.0440	-0.0000j	1.0440	-0.0000
<b>AR.2</b>	0.9467	-0.4502j	1.0483	-0.0707
<b>AR.3</b>	0.9467	+0.4502j	1.0483	0.0707
<b>AR.4</b>	-29.3860	-0.0000j	29.3860	-0.5000
<b>MA.1</b>	0.8781	-0.4785j	1.0000	-0.0794
<b>MA.2</b>	0.8781	+0.4785j	1.0000	0.0794
<b>MA.3</b>	-0.9745	-0.7755j	1.2454	-0.3930
<b>MA.4</b>	-0.9745	+0.7755j	1.2454	0.3930


```
In [168]: r2_score(logged_df_minus_roll_mean1, ar.predict())
```

```
Out[168]: 0.9930234354656668
```

- This means that 99.3 percent of the variation in the y data is due to variation in the x data
- This might indicate overfitting, but we chose our params from a stationary time series ACF and PACF.

-Future work: investigate more tweaks to the model

```
In [169]: #plot of ARMA model  
fig = ar.plot_predict()
```



**Change the params, maybe it will affect  $r^2$**

In [170]: *# Try p = 4 and q = 2*

```
# Instantiate & fit model with statsmodels
#p = num lags - ACF
p = 4

# q = lagged forecast errors - PACF
q = 2

# Fitting ARMA model and summary
ar = ARMA(logged_df_minus_roll_mean1,(p, q)).fit()
ar.summary()
```

Out[170]:

ARMA Model Results

<b>Dep. Variable:</b>	y	<b>No. Observations:</b>	255
<b>Model:</b>	ARMA(4, 2)	<b>Log Likelihood</b>	1189.035
<b>Method:</b>	csmle	<b>S.D. of innovations</b>	0.002
<b>Date:</b>	Thu, 29 Apr 2021	<b>AIC</b>	-2362.069
<b>Time:</b>	14:17:08	<b>BIC</b>	-2333.739
<b>Sample:</b>	02-01-1997	<b>HQIC</b>	-2350.674
	- 04-01-2018		

	coef	std err	z	P> z	[0.025	0.975]
const	0.0158	0.010	1.597	0.112	-0.004	0.035
ar.L1.y	1.1351	0.123	9.251	0.000	0.895	1.376
ar.L2.y	-0.3223	0.262	-1.231	0.219	-0.835	0.191
ar.L3.y	0.4592	0.284	1.619	0.107	-0.097	1.015
ar.L4.y	-0.3089	0.140	-2.209	0.028	-0.583	-0.035
ma.L1.y	1.0760	0.120	8.988	0.000	0.841	1.311
ma.L2.y	0.6758	0.076	8.867	0.000	0.526	0.825

Roots


	Real	Imaginary	Modulus	Frequency
AR.1	-0.5156	-1.3517j	1.4467	-0.3080
AR.2	-0.5156	+1.3517j	1.4467	0.3080
AR.3	1.0638	-0.0000j	1.0638	-0.0000
AR.4	1.4542	-0.0000j	1.4542	-0.0000
MA.1	-0.7961	-0.9197j	1.2164	-0.3635
MA.2	-0.7961	+0.9197j	1.2164	0.3635

```
In [171]: #Slightly lower r^2...hmm
r2_score(logged_df_minus_roll_mean1, ar.predict())
```

Out[171]: 0.9925556183138879

## Forecasting

```
In [172]: #plot of ARMA model
fig, ax = plt.subplots()
ax = logged_df_minus_roll_mean1.plot(ax=ax)
fig = ar.plot_predict('2018-05-01', '2021-04-01', dynamic=True, ax=ax, plot_insample=False)
plt.show()
```



Lower prices might be good indication to buy.

```
In [173]: #Future work, try SARIMAX prediction
# Need to install modules properly for SARIMAX to work.
```